Homework 5

HX

Table of contents

Question 1																				2	2
Question 2																				10)
Question 3															 					16	3

Link to the Github repository

Due: Wed, Apr 19, 2023 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

- 1. This assignment requires you to only upload a PDF file on Canvas
- 2. Don't collapse any code cells before submitting.
- 3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the author information in the frontmatter before submitting your assignment

In this assignment, we will explore decision trees, support vector machines and neural networks for classification and regression. The assignment is designed to test your ability to fit and analyze these models with different configurations and compare their performance.

We will need the following packages:

```
rm(list=ls())
packages <- c(
   "dplyr",
   "readr",</pre>
```

```
"tidyr",
   "purrr",
   "broom",
   "magrittr",
   "corrplot",
   "caret",
   "rpart",
   "rpart.plot",
   "e1071",
   "torch",
   "luz"
)

#renv::install(packages)
sapply(packages, require, character.only=T)
```

Question 1



Prediction of Median House prices

1.1 (2.5 points)

The data folder contains the housing.csv dataset which contains housing prices in California from the 1990 California census. The objective is to predict the median house price for California districts based on various features.

Read the data file as a tibble in R. Preprocess the data such that:

- 1. the variables are of the right data type, e.g., categorical variables are encoded as factors
- 2. all column names to lower case for consistency
- 3. Any observations with missing values are dropped

```
path <- "/Users/alex/Desktop/SP2023/STAT380/hw-5/data/housing.csv"

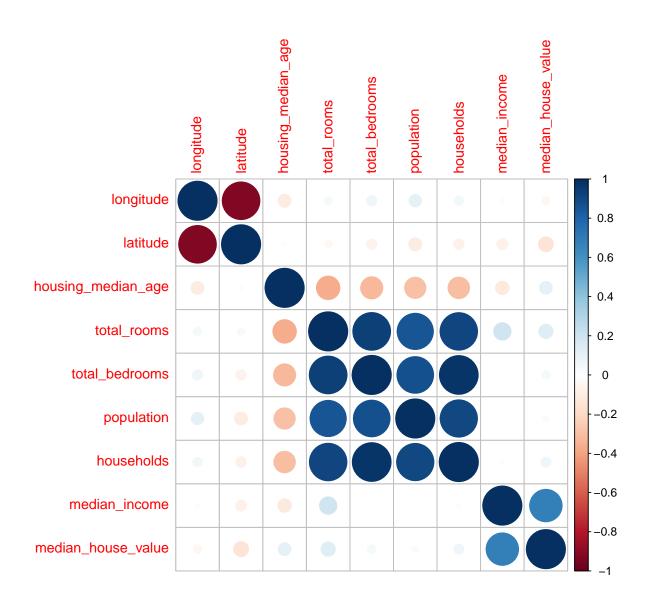
df <- read.csv(path)
  df$ocean_proximity<-as.factor(df$ocean_proximity)
  str(df)</pre>
```

```
20640 obs. of 10 variables:
'data.frame':
$ longitude
                    : num -122 -122 -122 -122 -122 ...
$ latitude
                     : num 37.9 37.9 37.9 37.9 37.9 ...
$ housing_median_age: num 41 21 52 52 52 52 52 52 52 42 52 ...
$ total rooms
                            880 7099 1467 1274 1627 ...
                   : num
$ total_bedrooms
                            129 1106 190 235 280 ...
                    : num
$ population
                    : num
                            322 2401 496 558 565 ...
$ households
                     : num 126 1138 177 219 259 ...
$ median income
                    : num 8.33 8.3 7.26 5.64 3.85 ...
$ median_house_value: num 452600 358500 352100 341300 342200 ...
$ ocean_proximity : Factor w/ 5 levels "<1H OCEAN", "INLAND", ...: 4 4 4 4 4 4 4 4 4 4 ...
  names(df) <- tolower(names(df))</pre>
  df <- na.omit(df)</pre>
```

1.2 (2.5 points)

Visualize the correlation matrix of all numeric columns in df using corrplot()

```
dfnum <- df%>%
  select(!ocean_proximity)
dfcor <- cor(dfnum)
corrplot(dfcor)</pre>
```



1.3 (5 points)

Split the data df into df_train and df_split using test_ind in the code below:

```
set.seed(42)
test_ind <- sample(</pre>
```

```
1:nrow(df),
floor( nrow(df)/10 ),
replace=FALSE
)

df_train <- df[-test_ind, ]
df_test <- df[test_ind, ]</pre>
```

1.4 (5 points)

Fit a linear regression model to predict the median_house_value:

- latitude
- longitude
- housing_median_age
- total_rooms
- total_bedrooms
- population
- median_income
- ocean_proximity

Interpret the coefficients and summarize your results.

```
lm_fit <- lm(median_house_value ~ latitude + longitude + housing_median_age +</pre>
                  total_rooms + total_bedrooms + population + median_income +
                  ocean_proximity, data = df_train)
  summary(lm_fit)
Call:
lm(formula = median_house_value ~ latitude + longitude + housing_median_age +
    total_rooms + total_bedrooms + population + median_income +
    ocean_proximity, data = df_train)
Residuals:
    Min
             1Q Median
                             3Q
                                    Max
-559024 -42322 -10389
                          28743 710215
Coefficients:
                            Estimate Std. Error t value Pr(>|t|)
```

```
(Intercept)
                         -2.273e+06 9.138e+04 -24.873 < 2e-16 ***
latitude
                         -2.539e+04 1.047e+03 -24.244 < 2e-16 ***
                         -2.681e+04 1.060e+03 -25.305 < 2e-16 ***
longitude
housing_median_age
                          1.074e+03 4.616e+01 23.261 < 2e-16 ***
total rooms
                         -6.159e+00 8.431e-01 -7.306 2.87e-13 ***
total_bedrooms
                          1.353e+02 4.254e+00 31.804 < 2e-16 ***
population
                         -3.413e+01 9.838e-01 -34.694
                                                       < 2e-16 ***
median_income
                          3.936e+04 3.573e+02 110.154 < 2e-16 ***
ocean_proximityINLAND
                         -4.018e+04 1.836e+03 -21.891
                                                       < 2e-16 ***
ocean_proximityISLAND
                          1.324e+05 3.442e+04
                                                 3.847
                                                        0.00012 ***
                         -2.522e+03 2.022e+03 -1.247
ocean_proximityNEAR BAY
                                                        0.21226
ocean_proximityNEAR OCEAN 4.349e+03 1.658e+03
                                                 2.622 0.00875 **
               0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Signif. codes:
Residual standard error: 68780 on 18378 degrees of freedom
Multiple R-squared: 0.643, Adjusted R-squared: 0.6428
F-statistic: 3009 on 11 and 18378 DF, p-value: < 2.2e-16
```

1.5 (5 points)

Complete the rmse function for computing the Root Mean-Squared Error between the true y and the predicted yhat, and use it to compute the RMSE for the regression model on df_test

```
rmse <- function(y, yhat) {
    sqrt(mean((y - yhat)^2))
}

lm_predictions <- predict(lm_fit, newdata = df_test)
lm_rmse <- rmse(df_test$median_house_value, lm_predictions)
lm_rmse

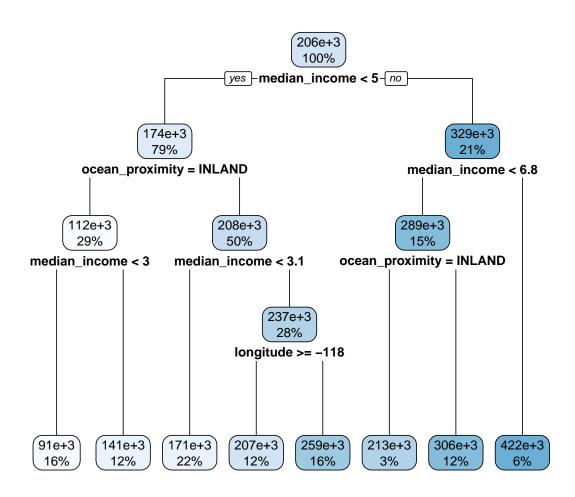
[1] 68339.82</pre>
```

1.6 (5 points)

Fit a decision tree model to predict the median_house_value using the same predictors as in 1.4. Use the rpart() function.

Visualize the decision tree using the rpart.plot() function.

```
rpart.plot(rpart_fit)
```



Report the root mean squared error on the test set.

```
rpart_predictions <- predict(rpart_fit, newdata = df_test)
rpart_rmse <- rmse(df_test$median_house_value, rpart_predictions)
rpart_rmse</pre>
[1] 75876.87
```

1.7 (5 points)

Fit a support vector machine model to predict the median_house_value using the same predictors as in 1.4. Use the svm() function and use any kernel of your choice. Report the root mean squared error on the test set.

[1] 56678.84

1.8 (25 points)

Initialize a neural network model architecture:

```
NNet <- nn_module(
  initialize = function(p, q1, q2, q3) {
    self$fc1 <- nn_linear(p, q1)
    self$fc2 <- nn_linear(q1, q2)
    self$fc3 <- nn_linear(q2, q3)
    self$relu <- nn_relu()
    self$dropout <- nn_dropout(p = 0.5)
},
forward = function(x) {
    x %>%
```

```
torch$nn$functional$linear(self$fc1, .) %>%
    self$relu(.) %>%
    self$dropout(.) %>%
    torch$nn$functional$linear(self$fc2, .) %>%
    self$relu(.) %>%
    self$dropout(.) %>%
    torch$nn$functional$linear(self$fc3, .)
}
```

Fit a neural network model to predict the median_house_value using the same predictors as in 1.4. Use the model.matrix function to create the covariate matrix and luz package for fitting the network with 32, 16, 8 nodes in each of the three hidden layers.

Plot the results of the training and validation loss and accuracy.

```
... # Insert your code here
```

Report the root mean squared error on the test set.

```
nnet_predictions <- ... # Insert your code here</pre>
```

⚠ Warning

Remember to use the as_array() function to convert the predictions to a vector of numbers before computing the RMSE with rmse()

1.9 (5 points)

Summarize your results in a table comparing the RMSE for the different models. Which model performed best? Why do you think that is?

Lower values of RMSE indicate better fit. Therefore, SVM is better.

```
lm_rmse <- rmse(df_test$median_house_value, lm_predictions)</pre>
  rpart_rmse <- rmse(df_test$median_house_value, rpart_predictions)</pre>
  svm_rmse <- rmse(df_test$median_house_value, svm_predictions)</pre>
  results <- data.frame(
    Model = c("Linear Regression", "Decision Tree", "Support Vector Machine"),
    RMSE = c(lm_rmse, rpart_rmse, svm_rmse)
  print(results)
                    Model
                              RMSE
1
       Linear Regression 68339.82
2
           Decision Tree 75876.87
3 Support Vector Machine 56678.84
```

Question 2



50 points

Spam email classification

The data folder contains the spam.csv dataset. This dataset contains features extracted from a collection of spam and non-spam emails. The objective is to classify the emails as spam or non-spam.

2.1 (2.5 points)

Read the data file as a tibble in R. Preprocess the data such that:

- 1. the variables are of the right data type, e.g., categorical variables are encoded as factors
- 2. all column names to lower case for consistency
- 3. Any observations with missing values are dropped

```
df2 <- read.csv("/Users/alex/Desktop/SP2023/STAT380/hw-5/data/spambase.csv")
invisible(lapply(df2,as.numeric))
names(df2) <- tolower(names(df2))
df2 <- na.omit(df2)</pre>
```

2.2 (2.5 points)

Split the data df into df_train and df_split using test_ind in the code below:

```
set.seed(42)
test_ind <- sample(
   1:nrow(df2),
   floor( nrow(df2)/10 ),
   replace=FALSE
)

df2_train <- df2[-test_ind, ]
df2_test <- df2[test_ind, ]</pre>
```

Complete the overview function which returns a data frame with the following columns: accuracy, error, false positive rate, true positive rate, between the true true_class and the predicted pred_class for any classification model.

```
overview <- function(pred_class, true_class) {
  accuracy <- sum(pred_class == true_class) / length(true_class)
  error <- 1 - accuracy
  true_positives <- sum(pred_class == true_class & true_class == 1)
  true_negatives <- sum(pred_class == true_class & true_class == 0)</pre>
```

```
false_positives <- sum(pred_class != true_class & true_class == 0)
false_negatives <- sum(pred_class != true_class & true_class == 1)
true_positive_rate <- true_positives / sum(true_class == 1)
false_positive_rate <- false_positives / sum(true_class == 0)
return(
    data.frame(
    accuracy = accuracy,
    error = error,
    true_positive_rate = true_positive_rate,
    false_positive_rate = false_positive_rate
    )
)
)</pre>
```

2.3 (5 points)

Fit a logistic regression model to predict the spam variable using the remaining predictors. Report the prediction accuracy on the test set.

```
glm_fit <- glm(spam ~ ., data = df2_train, family = "binomial")</pre>
```

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```
glm_probs <- predict(glm_fit, newdata = df2_test, type = "response")
glm_classes <- ifelse(glm_probs > 0.5, 1, 0)
accuracy <- mean(glm_classes == df2_test$spam)
accuracy</pre>
```

[1] 0.923913

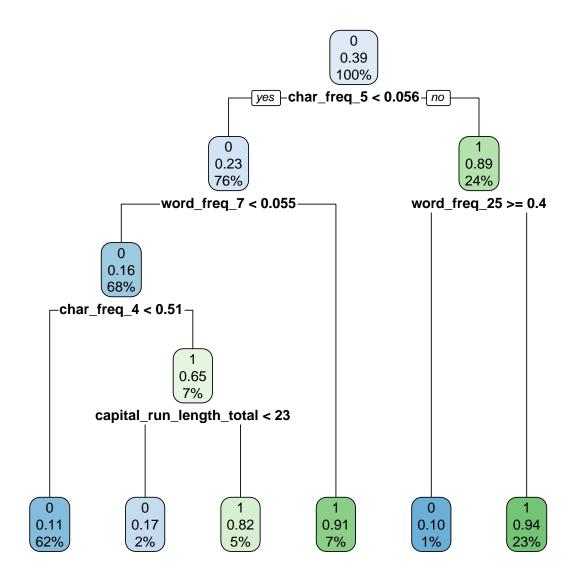
2.4 (5 points)

Fit a decision tree model to predict the spam variable using the remaining predictors. Use the rpart() function and set the method argument to "class".

```
rpart_fit2 <- rpart(spam ~ ., data = df2_train, method = "class")
rpart_probs2 <- predict(rpart_fit2, newdata = df2_test, type = "prob")
rpart_classes2 <- ifelse(rpart_probs2[, 2] > 0.5, 1, 0)
```

Visualize the decision tree using the rpart.plot() function.

```
rpart.plot(rpart_fit2)
```



Report the prediction accuracy on the test set.

```
rpart_classes2 <- predict(rpart_fit2, newdata = df2_test, type = "class")</pre>
rpart_acc2 <- mean(rpart_classes2 == df2_test$spam)</pre>
rpart_acc2
```

2.5 (5 points)

[1] 0.8782609

Fit a support vector machine model to predict the spam variable using the remaining predictors. Use the sym() function and use any kernel of your choice. Remember to set the type argument to "C-classification" if you haven't already converted spam to be of type factor.

```
if (!is.factor(df2_train$spam)) {
  df2_train$spam <- as.factor(df2_train$spam)</pre>
svm_fit2 <- svm(spam ~ ., data = df2_train, kernel = "radial")</pre>
```

Report the prediction accuracy on the test set.

```
svm_classes2 <- predict(svm_fit2, newdata = df2_test)</pre>
svm acc2 <- mean(svm classes2 == df2 test$spam)</pre>
svm_acc2
```

[1] 0.923913

2.6 (25 points)

Using the same neural network architecture as in 1.9, fit a neural network model to predict the spam variable using the remaining predictors.

△ Classification vs. Regression

Note that the neural network in Q 1.9 was a regression model. You will need to modify the neural network architecture to be a classification model by changing the output layer to have a single node with a sigmoid activation function.

Use the model.matrix function to create the covariate matrix and luz package for fitting the network with 32,16,8 nodes in each of the three hidden layers.

```
X_{train} \leftarrow model.matrix(spam \sim ., data = df2_train)[,-1] # Remove the intercept
Y_train <- df2_train$spam
X_test <- model.matrix(spam ~ ., data = df2_test)[,-1]</pre>
Y_test <- df2_test$spam
normalize <- function(x) {</pre>
  return((x - min(x)) / (max(x) - min(x)))
}
X_train <- as.data.frame(apply(X_train, 2, normalize))</pre>
X_test <- as.data.frame(apply(X_test, 2, normalize))</pre>
nn_model <- luz::net() %>%
  input(57) %>% # 57 predictors
  dense(32) %>%
  relu() %>%
  dense(16) %>%
  relu() %>%
  dense(8) %>%
  relu() %>%
  dense(1) %>%
  sigmoid()
set.seed(42)
nn fit <- nn model %>%
  luz::optimizer("adam", lr = 0.001) %>%
  luz::loss("binary_crossentropy") %>%
  luz::train(
    X_train,
    Y_train,
    batch_size = 32,
    epochs = 100,
    validation_data = list(X_test, Y_test)
nn_probs <- predict(nn_fit, X_test)</pre>
nn_classes <- ifelse(nn_probs > 0.5, 1, 0)
```

```
nn_accuracy <- mean(nn_classes == Y_test)</pre>
nn_accuracy
```

2.7 (5 points)

Summarize your results in a table comparing the accuracy metrics for the different models.

Accuracy of Support vector machine and log regression are the same.

```
results2 <- data.frame(Model = c('Logistic Regression', 'Decision Tree', 'Support Vector M
                      Accuracy = c(accuracy,rpart_acc2,svm_acc2))
results2
```

```
Model
                          Accuracy
1
     Logistic Regression 0.9239130
           Decision Tree 0.8782609
3 Support Vector Machine 0.9239130
```

If you were to choose a model to classify spam emails, which model would you choose? Think about the context of the problem and the cost of false positives and false negatives.

Question 3



🥊 60 points

Three spirals classification

To better illustrate the power of depth in neural networks, we will use a toy dataset called the "Three Spirals" data. This dataset consists of two intertwined spirals, making it challenging for shallow models to classify the data accurately.



⚠ This is a multi-class classification problem

The dataset can be generated using the provided R code below:

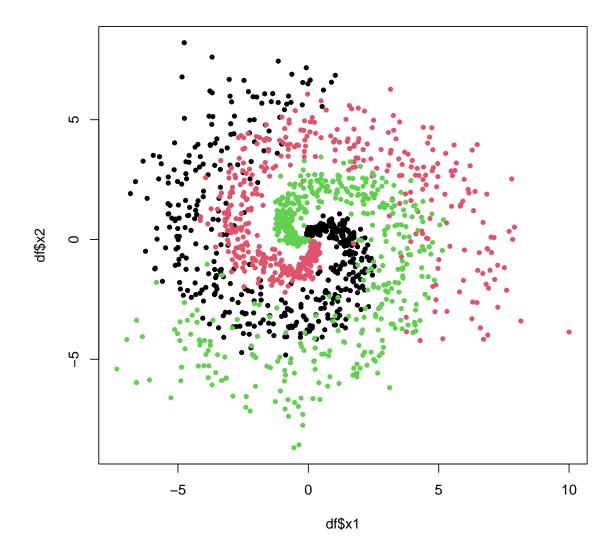
```
generate_three_spirals <- function(){</pre>
 set.seed(42)
 n <- 500
 noise <- 0.2
 t <- (1:n) / n * 2 * pi
  x1 <- c(
     t * (sin(t) + rnorm(n, 0, noise)),
      t * (sin(t + 2 * pi/3) + rnorm(n, 0, noise)),
      t * (sin(t + 4 * pi/3) + rnorm(n, 0, noise))
 x2 <- c(
     t * (cos(t) + rnorm(n, 0, noise)),
      t * (cos(t + 2 * pi/3) + rnorm(n, 0, noise)),
     t * (cos(t + 4 * pi/3) + rnorm(n, 0, noise))
 y <- as.factor(
   c(
     rep(0, n),
     rep(1, n),
     rep(2, n)
  )
 return(tibble(x1=x1, x2=x2, y=y))
```

3.1 (5 points)

Generate the three spirals dataset using the code above. Plot x_1 vs x_2 and use the y variable to color the points.

```
df <- generate_three_spirals()

plot(
   df$x1, df$x2,
   col = df$y,
   pch = 20
)</pre>
```



Define a grid of 100 points from -10 to 10 in both x_1 and x_2 using the expand.grid(). Save it as a tibble called df_test.

```
library(tibble)
grid <- expand.grid(
    x1 = seq(-10, 10, length.out = 100),
    x2 = seq(-10, 10, length.out = 100)
)</pre>
```

```
df_test <- as_tibble(grid)</pre>
```

3.2 (10 points)

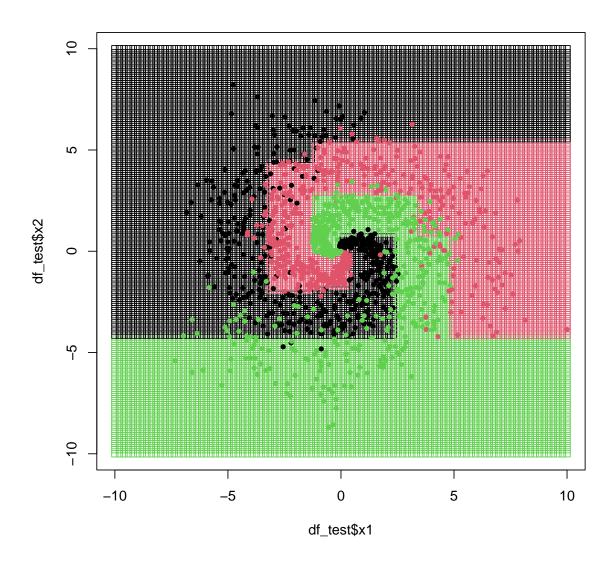
Fit a classification tree model to predict the y variable using the x1 and x2 predictors, and plot the decision boundary.

```
rpart_fit <- rpart(y ~ x1 + x2, data = df, method = "class")
rpart_classes <- predict(rpart_fit, newdata = df_test, type = "class")</pre>
```

Plot the decision boundary using the following function:

```
plot_decision_boundary <- function(predictions){
  plot(
    df_test$x1, df_test$x2,
    col = predictions,
    pch = 0
)
  points(
    df$x1, df$x2,
    col = df$y,
    pch = 20
)
}

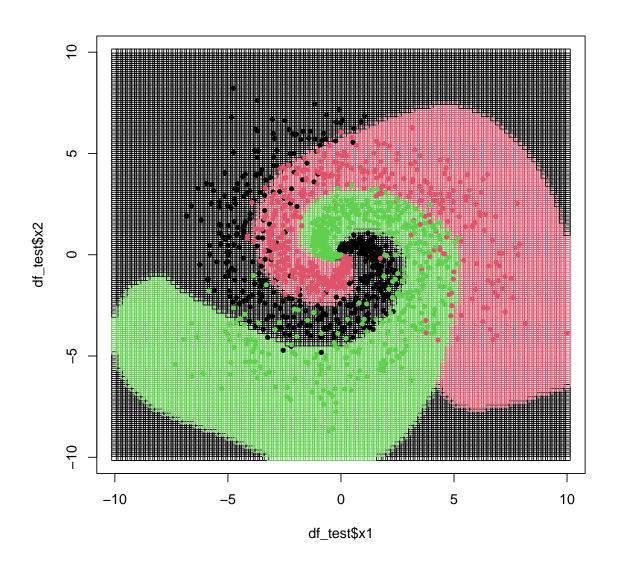
plot_decision_boundary(rpart_classes)</pre>
```



3.3 (10 points)

Fit a support vector machine model to predict the y variable using the x1 and x2 predictors. Use the svm() function and use any kernel of your choice. Remember to set the type argument to "C-classification" if you haven't converted y to be of type factor.

svm_fit <- svm(y ~ x1 + x2, data = df, kernel = "radial", type = "C-classificat
svm_classes <- predict(svm_fit, newdata = df_test)
plot_decision_boundary(svm_classes)</pre>



⚠ Instructions

For the next questions, you will need to fit a series of neural networks. In all cases, you can:

- set the number of units in each hidden layer to 10
- set the output dimension o to 3 (remember this is multinomial classification)
- use the appropriate loss function for the problem (not nn_bce_loss)
- set the number of epochs to 50
- fit the model using the luz package

You can use any optimizer of your choice, but you will need to tune the learning rate for each problem.

3.4 (10 points)

Fit a neural network with 1 hidden layer to predict the y variable using the x1 and x2 predictors.

```
library(torch)
library(luz)
NN1 <- nn_module(</pre>
  initialize = function(p, q1, o){
     self$hidden1 <- nn_layer(q1, p)</pre>
     self$output <- nn_layer(o, q1)</pre>
     self$activation <- nn_sigmoid()</pre>
  },
  forward = function(x){
    x %>%
       self$hidden1() %>%
       self$activation() %>%
       self$output()
  }
)
\mathtt{fit}\_1 \; \leftarrow \; \mathtt{NN1} \; \% \gt \%
  setup(
     input_shape = ncol(df) - 1,
     output shape = 3,
    hidden_shapes = list(q1 = 10)
  ) %>%
  set_hparams(
```

```
epochs = 50,
    learning_rate = 0.001
  ) %>%
  set_optimizers(
    "adam",
    lr = 0.001
  ) %>%
  fit(
    data = list(
      df %>% select(x1, x2) %>% as.matrix,
      df$y %>% as.integer
    ),
    dataloader_options = list(batch_size = 128, shuffle = TRUE),
    verbose = FALSE
  )
test_matrix <- df_test %>% select(x1, x2) %>% as.matrix
fit_1_predictions <- predict(fit_1, test_matrix) %>%
  argmax(2) %>%
  as.integer()
plot_decision_boundary(fit_1_predictions)
```

In order to generate the class predictions, you will need to use the predict() function as follows

```
test_matrix <- df_test %>% select(x1, x2) %>% as.matrix
library(ramify)
fit_1_predictions <- predict(fit_1, test_matrix) %>%
    argmax(2) %>%
    as.integer()
```

Plot the results using the plot_decision_boundary() function.

3.5 (10 points)

Fit a neural network with 0 hidden layers to predict the y variable using the x1 and x2 predictors.

```
NNO <- nn_module(
  initialize = function(p, o){
    self$output <- nn_layer(o)</pre>
  },
  forward = function(x){
    x %>%
    self$output()
  }
)
fit_0 <- NNO %>%
  setup(input_size = 2, output_size = 3) %>%
  set_hparams(units = NULL, loss = nn_ce_loss, epochs = 50) %>%
  set_optimizers(adam_optimizer(learning_rate = 0.01)) %>%
  fit(
    data = list(
      df %>% select(x1, x2) %>% as.matrix,
      df$y %>% as.integer
    ),
    dataloader_options = list(batch_size = 128, shuffle = TRUE),
    verbose = FALSE
  )
```

Plot the results using the plot_decision_boundary() function.

```
fit_0_predictions <- predict(fit_0, test_matrix) %>%
   argmax(2) %>%
   as.integer()

plot_decision_boundary(fit_0_predictions)
```

3.6 (10 points)

Fit a neural network with 3 hidden layers to predict the y variable using the x1 and x2 predictors.

```
NN2 <- nn_module(
  initialize = function(p, q1, q2, q3, o){
    self$hidden1 <- nn_linear(q1, p)
    self$hidden2 <- nn_linear(q2, q1)</pre>
```

```
self$hidden3 <- nn_linear(q3, q2)</pre>
    self$output <- nn_linear(o, q3)</pre>
    self$activation <- nn_functional$relu</pre>
  },
  forward = function(x){
    x %>%
      self$hidden1() %>%
      self$activation() %>%
      self$hidden2() %>%
      self$activation() %>%
      self$hidden3() %>%
      self$activation() %>%
      self$output()
  }
)
fit_2 <- NN2 %>%
  setup(
    input_size = 2,
    output_size = 3
  ) %>%
  set hparams(
    hidden_sizes = c(10, 10, 10),
    activation_function = "relu",
    learning_rate = 0.01
  ) %>%
  set_optimizers(
    optimizer = "adam"
  ) %>%
  fit(
    X_train = df %>% select(x1, x2) %>% as.matrix(),
    Y_train = df$y %>% as.factor() %>% as.numeric() %>% as.matrix(),
    n_{epoch} = 50,
    batch_size = 128,
    shuffle = TRUE,
    verbose = FALSE
  )
```

Plot the results using the plot_decision_boundary() function.

```
test_matrix <- df_test %>% select(x1, x2) %>% as.matrix()
fit_2_predictions <- predict(fit_2, test_matrix) %>%
    argmax(2) %>%
    as.integer()

plot_decision_boundary(df_test, fit_2_predictions)
```

3.7 (5 points)

What are the differences between the models? How do the decision boundaries change as the number of hidden layers increases?

The three models are different in terms of the number of hidden layers used to make predictions. As the complexity of the decision boundary increases, the model becomes better able to classify points accurately.

i Session Information Print your R session information using the following command sessionInfo() R version 4.2.3 (2023-03-15) Platform: x86_64-apple-darwin17.0 (64-bit) Running under: macOS Big Sur ... 10.16 Matrix products: default /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRblas.0.dylib BLAS: LAPACK: /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRlapack.dylib locale: [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8 attached base packages: graphics grDevices datasets utils [1] stats methods base other attached packages: [1] tibble_3.2.1 luz_0.4.0 torch_0.10.0 e1071 1.7-13 [5] rpart.plot_3.1.1 rpart_4.1.19 $lattice_0.20-45$ caret_6.0-94 [9] ggplot2_3.4.2 corrplot_0.92 magrittr_2.0.3 broom_1.0.4 [13] purrr_1.0.1 tidyr_1.3.0 readr_2.1.4 dplyr_1.1.1 loaded via a namespace (and not attached): [1] nlme_3.1-162 $fs_1.6.1$ lubridate_1.9.2 [4] bit64_4.0.5 progress_1.2.2 tools_4.2.3 [7] backports_1.4.1 utf8_1.2.3 R6_2.5.1 [10] colorspace_2.1-0 nnet_7.3-18 withr_2.5.0 [13] tidyselect_1.2.0 prettyunits_1.1.1 processx_3.8.0 [16] bit_4.0.5 compiler_4.2.3 cli_3.6.1 [19] scales_1.2.1 callr_3.7.3 proxy_0.4-27 [22] stringr 1.5.0 digest_0.6.31 rmarkdown_2.21 [25] coro_1.0.3 pkgconfig_2.0.3 htmltools_0.5.5

rlang_1.1.0

Rcpp_1.0.10

jsonlite_1.8.4

lifecycle_1.0.3

fastmap_1.1.1

fansi_1.0.4

generics_0.1.3

[28] parallelly_1.35.0

[34] ModelMetrics_1.2.2.2 Matrix_1.5-3

[31] rstudioapi_0.14

[37] munsell_0.5.0

F407	707 4 40 0	
[40] stringi_1.7.12	pROC_1.18.0	yaml_2.3.7
[43] MASS_7.3-58.2	plyr_1.8.8	recipes_1.0.5
[46] grid_4.2.3	parallel_4.2.3	listenv_0.9.0
[49] crayon_1.5.2	splines_4.2.3	hms_1.1.3
[52] zeallot_0.1.0	knitr_1.42	ps_1.7.4
[55] pillar_1.9.0	future.apply_1.10.0	reshape2_1.4.4
[58] codetools_0.2-19	stats4_4.2.3	glue_1.6.2
[61] evaluate_0.20	data.table_1.14.8	renv_0.16.0-53
[64] vctrs_0.6.1	tzdb_0.3.0	foreach_1.5.2
[67] gtable_0.3.3	future_1.32.0	xfun_0.38
[70] gower_1.0.1	prodlim_2023.03.31	class_7.3-21
[73] survival_3.5-3	timeDate_4022.108	iterators_1.0.14
[76] hardhat_1.3.0	lava_1.7.2.1	timechange_0.2.0
[79] globals_0.16.2	ipred_0.9-14	